

7 Traffic prediction

7. (15 points) You were just hired by Ways, and they ask you to use machine-learning methods to predict the traffic at intersections in a city.

(a) The first thing you have to decide is how to encode the inputs. The first city you analyze data for is designed on a grid, where all streets either run due north/south or due east/west. You know the names of all the streets in the city in advance. You have measurements of the traffic at many intersections at many times of day, and you have to encode each measurement using a vector of feature values.

For each of the features below, explain how best to encode it. Your encoding should be able to take into account both the fact that traffic at different points on the same road tends to be correlated and the fact that some whole parts of a city might be busier than others.

- i. Name of street that runs north/south
- ii. Name of street that runs east/west

- iii. Latitude (positive real)
- iv. Longitude (positive real)

- v. Day of week (an integer from 1 to 7)

- vi. Time of day (a real number between 0 and 24)

Name: _____

- (b) You do such a good job in that city, you are assigned to work on a new city that is very old, and although there are still only two streets involved in an intersection, the streets wind all around and don't run in a consistent direction, so it's not clear which order to put them in. So, a training example for an intersection might be arbitrarily described as one of

("1st street", "park street", 34.4, 54.2, 02139, 22.35, Friday)

("park street", "1st street", 34.4, 54.2, 02139, 22.35, Friday)

You can assume that for any two streets, there is only a single intersection between them. How would you change your encoding of the first two features to deal with this, or would you leave it as is? Explain your answer.

- (c) Your goal is to predict how much traffic there is at a given intersection at a given time, so now you have to think about how to encode the output. We'll just consider a single direction of a single road. The raw data that you have contains either a positive number or `None`. If it's a number, that is the average speed in miles per hour of the cars that went through this road in the past 10 minutes. If it is `None`, then no cars went through in the last 10 minutes.

Design an output encoding for this data. Specify how many dimensions it has, and precisely what transformation you would do on the raw output values in order to compute your encoding.

Name: _____

- (d) You are using a neural network to predict traffic using the encoding you selected above. What activation function(s) would you use on your output units, and what loss function would you use? Explain your choices.